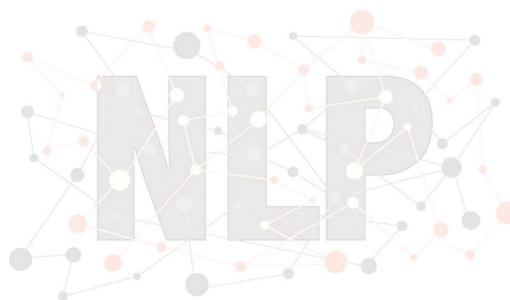


Unit 7

Applications of NLP

Natural Language Processing (NLP)
MDS 555



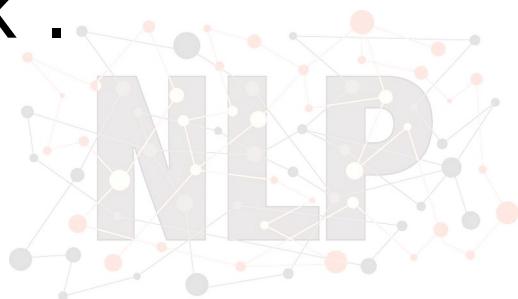
Objective

- N-Gram Language Model
- TF-IDF
 - Algorithm
 - Implementation – already done
- Vector Semantics and Embeddings



N-Gram Language Models

- Models that assign **probabilities** to sequences of words are called **language models** or LMs
- An n-gram is a sequence n-gram of n words
 - **2-gram** (**a bigram**) is a two-word sequence of words like “please turn”, “turn your”, or “your homework”
 - **3-gram** (**a trigram**) is a three-word sequence of words like “please turn your”, or “turn your homework”.



N-Grams

- The probability of a word w given some history h

$$P(w|h)$$

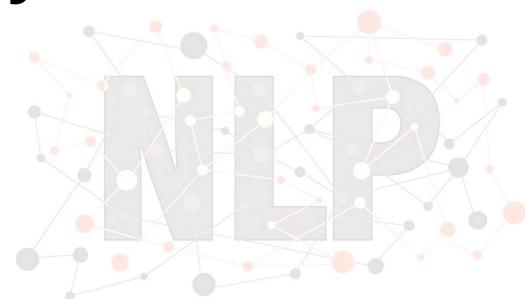
- Suppose the history h is “its water is so transparent that” and we want to know the probability that the next word is “the”

$$P(\text{the}|\text{its water is so transparent that})$$



N-Grams

- One way to estimate this probability is from relative frequency counts:
 - take a very large corpus,
 - count the number of times we see “**its water is so transparent that**”, and
 - count the number of times this is followed by “**the**”

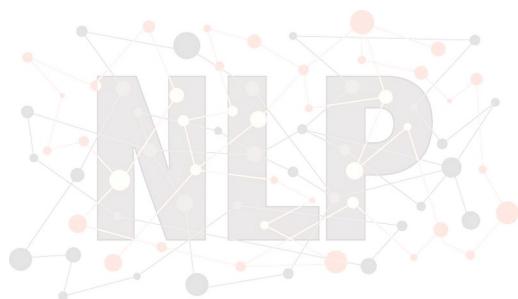


N-Grams

- This would be answering the question “Out of the times we saw the history h, how many times was it followed by the word w”, as follows

$$P(\text{the}|\text{its water is so transparent that}) = \frac{C(\text{its water is so transparent that the})}{C(\text{its water is so transparent that})}$$

- In many case counting this way will not results to the solution due to dynamic nature of the language
- It is not appropriate to count in larger corpus



N-Grams

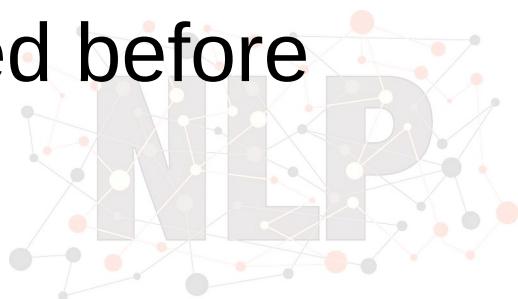
- **Chain Rule of Probability**
 - The chain rule shows the link between computing the joint probability of a sequence and computing the conditional probability of a word given previous words

$$\begin{aligned} P(X_1 \dots X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_{1:2}) \dots P(X_n|X_{1:n-1}) \\ &= \prod_{k=1}^n P(X_k|X_{1:k-1}) \end{aligned}$$



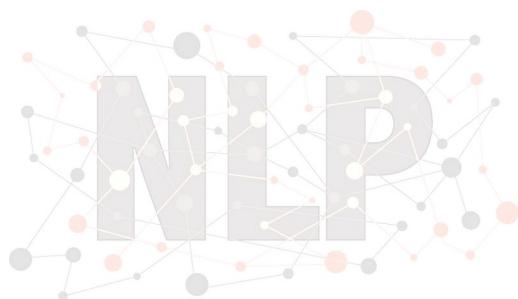
N-Grams

- We don't know any way to compute the exact probability of a word given a long sequence of preceding words,
 $P(w_n|w_{1:n-1})$.
- As we said above, we can't just estimate by counting the number of times every word occurs following every long string, because
 - language is creative and
 - any particular context might have never occurred before



Bigram

- The bigram model
 - Approximates the probability of a word given all the previous words $P(w_n|w_{1:n-1})$ by using only the conditional probability of the preceding word
 $P(w_n|w_{n-1})$
- In other words, instead of computing the probability
 $P(\text{the}|\text{Walden Pond's water is so transparent that})$
 - we approximate it with the probability
 $P(\text{the}|\text{that})$



Bigram

- When we **use a bigram model to predict the conditional probability** of the next word, we are thus making the following approximation:

$$P(w_n | w_{1:n-1}) \approx P(w_n | w_{n-1})$$

- The assumption that the probability of a word depends only on the previous word is called a **Markov assumption**



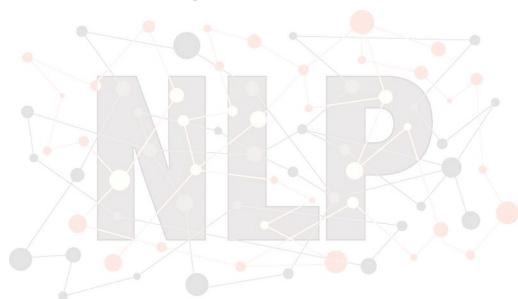
N-Gram Approximation

- **N** mean the n-gram size,
 - N = 2 means **bigrams** and N = 3 means **trigrams**
- Then we **approximate** the probability of a word given its entire context as follows:

$$P(w_n | w_{1:n-1}) \approx P(w_n | w_{n-N+1:n-1})$$

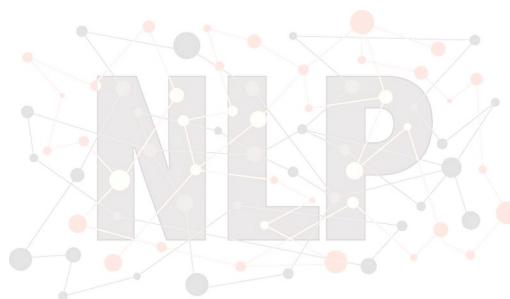
- Given the bigram assumption for the probability of an individual word, we can compute the probability of a complete word sequence by

$$P(w_{1:n}) \approx \prod_{k=1}^n P(w_k | w_{k-1})$$



Maximum likelihood estimation - MLE

- An intuitive way to estimate probabilities
- We get MLE estimate for the parameters of an n-gram model by **getting counts from a corpus, and normalizing the counts** so that they lie between 0 and 1



Maximum likelihood estimation - MLE

- An intuitive way to estimate probabilities
- We get MLE estimate for the parameters of an n-gram model by **getting counts from a corpus, and normalizing the counts** so that they lie between 0 and 1
 - To compute a particular bigram probability of a word w_n given a previous word w_{n-1} ,
 - we'll compute the count of the bigram $C(w_{n-1}w_n)$ and normalize by the sum of all the bigrams that share the same first word w_{n-1}

$$P(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n)}{\sum_w C(w_{n-1}w)}$$

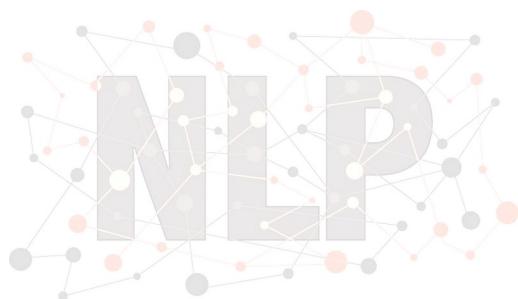


MLE

$$P(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n)}{\sum_w C(w_{n-1}w)}$$

- We can simplify this equation,
 - the sum of all bigram counts that start with a given word w_{n-1} must be equal to the unigram count for that word w_{n-1}

$$P(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n)}{C(w_{n-1})}$$



Bigram Counts

| | i | want | to | eat | chinese | food | lunch | spend |
|----------------|----------|-------------|-----------|------------|----------------|-------------|--------------|--------------|
| i | 5 | 827 | 0 | 9 | 0 | 0 | 0 | 2 |
| want | 2 | 0 | 608 | 1 | 6 | 6 | 5 | 1 |
| to | 2 | 0 | 4 | 686 | 2 | 0 | 6 | 211 |
| eat | 0 | 0 | 2 | 0 | 16 | 2 | 42 | 0 |
| chinese | 1 | 0 | 0 | 0 | 0 | 82 | 1 | 0 |
| food | 15 | 0 | 15 | 0 | 1 | 4 | 0 | 0 |
| lunch | 2 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| spend | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |

Figure 3.1 Bigram counts for eight of the words (out of $V = 1446$) in the Berkeley Restaurant Project corpus of 9332 sentences. Zero counts are in gray.



Bigram probabilities

| i | want | to | eat | chinese | food | lunch | spend |
|------|------|------|-----|---------|------|-------|-------|
| 2533 | 927 | 2417 | 746 | 158 | 1093 | 341 | 278 |

| | i | want | to | eat | chinese | food | lunch | spend |
|---------|---------|------|--------|--------|---------|--------|--------|---------|
| i | 0.002 | 0.33 | 0 | 0.0036 | 0 | 0 | 0 | 0.00079 |
| want | 0.0022 | 0 | 0.66 | 0.0011 | 0.0065 | 0.0065 | 0.0054 | 0.0011 |
| to | 0.00083 | 0 | 0.0017 | 0.28 | 0.00083 | 0 | 0.0025 | 0.087 |
| eat | 0 | 0 | 0.0027 | 0 | 0.021 | 0.0027 | 0.056 | 0 |
| chinese | 0.0063 | 0 | 0 | 0 | 0 | 0.52 | 0.0063 | 0 |
| food | 0.014 | 0 | 0.014 | 0 | 0.00092 | 0.0037 | 0 | 0 |
| lunch | 0.0059 | 0 | 0 | 0 | 0 | 0.0029 | 0 | 0 |
| spend | 0.0036 | 0 | 0.0036 | 0 | 0 | 0 | 0 | 0 |

Figure 3.2 Bigram probabilities for eight words in the Berkeley Restaurant Project corpus of 9332 sentences. Zero probabilities are in gray.



Use of bigram probability table

Here are a few other useful probabilities:

$$P(i | \langle s \rangle) = 0.25$$

$$P(\text{english} | \text{want}) = 0.0011$$

$$P(\text{food} | \text{english}) = 0.5$$

$$P(\langle /s \rangle | \text{food}) = 0.68$$

Now we can compute the probability of sentences like *I want English food* or *I want Chinese food* by simply multiplying the appropriate bigram probabilities together, as follows:

$$\begin{aligned} & P(\langle s \rangle \ i \ \text{want} \ \text{english} \ \text{food} \ \langle /s \rangle) \\ &= P(i | \langle s \rangle)P(\text{want} | i)P(\text{english} | \text{want}) \\ & \quad P(\text{food} | \text{english})P(\langle /s \rangle | \text{food}) \\ &= .25 \times .33 \times .0011 \times 0.5 \times 0.68 \\ &= .000031 \end{aligned}$$



Evaluating Language Model

- **Extrinsic evaluation**
 - End-to-End
 - The best way to evaluate the performance of a language model is to embed it in an application and measure **how much the application improves**
 - Example: Speech recognition
 - we can compare the performance of two language models by running the speech recognizer twice,
 - once with each language model, and seeing which gives the more accurate transcription.
 - It is **very expensive**



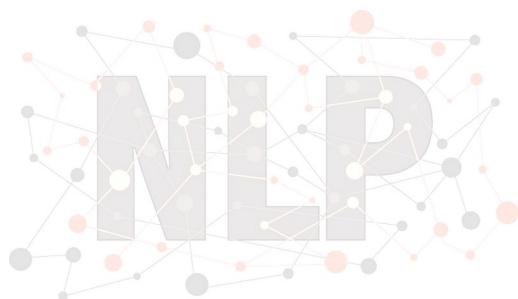
Evaluating Language Model

- Intrinsic evaluation
 - Measures the quality of a model independent of any application
 - We need a **test set**
 - We can then measure the quality of an n-gram model by its performance on some unseen data called the test set or test corpus



Limitation of Probabilistic Evaluation

- It's important not to let the test sentences into the training set
- If our test sentence is part of the training corpus, we will mistakenly assign it an artificially high probability when it occurs in the test set. We call this situation **training on the test set.**



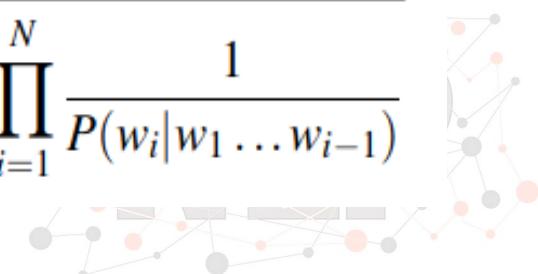
Perplexity (PPL)

- In practice we don't use raw probability as our metric for evaluating language models, but a variant called **perplexity**.
- The perplexity of a language model on a test set is the **inverse probability of the test set, normalized by the number of words**.

$$\begin{aligned}\text{perplexity}(W) &= P(w_1 w_2 \dots w_N)^{-\frac{1}{N}} \\ &= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}\end{aligned}$$

We can use the chain rule to expand the probability of W :

$$\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$



Self Study

- BOOK
 - Speech and Language Processing (3rd Edition)
Chapter 3 : N-gram Language Models



Text Vectorization



Text Vectorization

- **Text Vectorization** is the process of converting text into numerical representation.
- A technique for converting text into finite length vectors
 - Bag-of-Words
 - TF-IDF
 - CBOW
 - Skip Gram
 - Word2Vec
- Code the text into the numeric values



Bag of words

- We represent a text document as if it were a bag of words, that is, an **unordered set of words** with their **position ignored**, keeping only their frequency in the document.

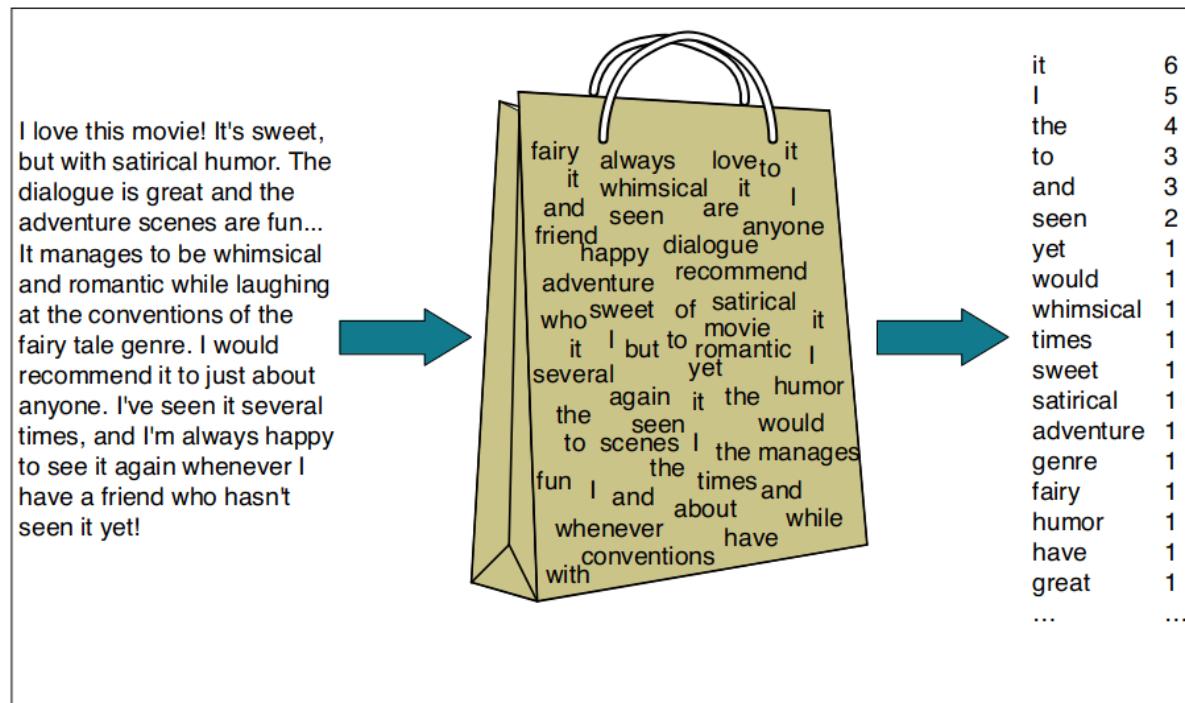
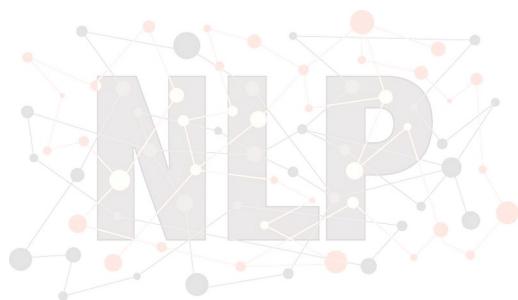


Figure 4.1 Intuition of the multinomial naive Bayes classifier applied to a movie review. The position of the words is ignored (the *bag-of-words* assumption) and we make use of the frequency of each word.



NLP Applications - Classifiers

- Self study
 - Text / Document Classification
 - Naive Bayes
 - **Self study:** how Naive Bayes can be used to build Language model
 - Chapter 4: Naive Bayes and Sentiment Classification
 - Evaluation: Precision, Recall and F-measure



Vector Semantics



Distributional Hypothesis

- Words that occur in **similar contexts tend to have similar meanings**. This link between similarity in how words are distributed and similarity in what they mean is called the **distributional hypothesis**.
 - The hypothesis was distributional hypothesis first formulated in the 1950s by linguists like Joos (1950), Harris (1954), and Firth (1957),
 - who noticed that
 - words which are synonyms (like oculist and eye-doctor) tended to occur in the same environment (e.g., near words like eye or examined) with the amount of meaning difference between two words “**corresponding roughly to the amount of difference in their environments**” (Harris, 1954, 1957).



Vector Semantics

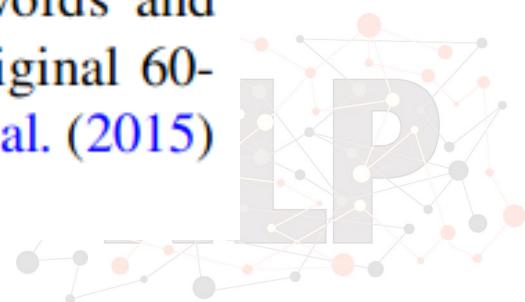
- Vector semantics is the standard way to **represent word meaning** in NLP, helping us model many of the aspects of word meaning
- The idea of vector semantics is to represent a word as a point in a **multidimensional semantic space** that is **derived from the distributions of word neighbors**.
- Vectors for representing words are called **embeddings**
 - The word “**embedding**” derives from its mathematical sense as a mapping from one space or structure to another, although the meaning has shifted



Embeddings



Figure 6.1 A two-dimensional (t-SNE) projection of embeddings for some words and phrases, showing that words with similar meanings are nearby in space. The original 60-dimensional embeddings were trained for sentiment analysis. Simplified from Li et al. (2015) with colors added for explanation.



Words and Vectors

- Vector or distributional models of meaning are generally based on a co-occurrence matrix, a way of representing how often words co-occur.
- Two popular matrices:
 - Document Dimension: the term-document matrix and
 - Word Dimension: the term-term matrix



Term-document matrix

- In a **term-document matrix**,
 - each row represents a word in the vocabulary and
 - each term-document matrix column represents a document from some collection of documents
 - The term-document matrix of Fig. 6.2 was first defined as part of the **vector space model of information retrieval** (Salton, 1971).

| | As You Like It | Twelfth Night | Julius Caesar | Henry V |
|--------|----------------|---------------|---------------|---------|
| battle | 1 | 0 | 7 | 13 |
| good | 114 | 80 | 62 | 89 |
| fool | 36 | 58 | 1 | 4 |
| wit | 20 | 15 | 2 | 3 |

Figure 6.2 The term-document matrix for four words in four Shakespeare plays. Each cell contains the number of times the (row) word occurs in the (column) document.

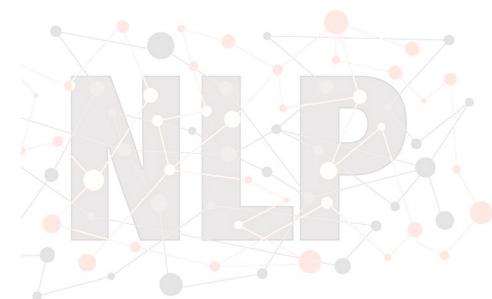


Term-document matrix

- The term-document matrix has
 - $|V|$ rows (one for each word type in the vocabulary) and
 - D columns (one for each document in the collection)

| | As You Like It | Twelfth Night | Julius Caesar | Henry V |
|--------|----------------|---------------|---------------|---------|
| battle | 1 | 0 | 7 | 13 |
| good | 114 | 80 | 62 | 89 |
| fool | 36 | 58 | 1 | 4 |
| wit | 20 | 15 | 2 | 3 |

Figure 6.3 The term-document matrix for four words in four Shakespeare plays. The red boxes show that each document is represented as a column vector of length four.



Visualizing the vectors

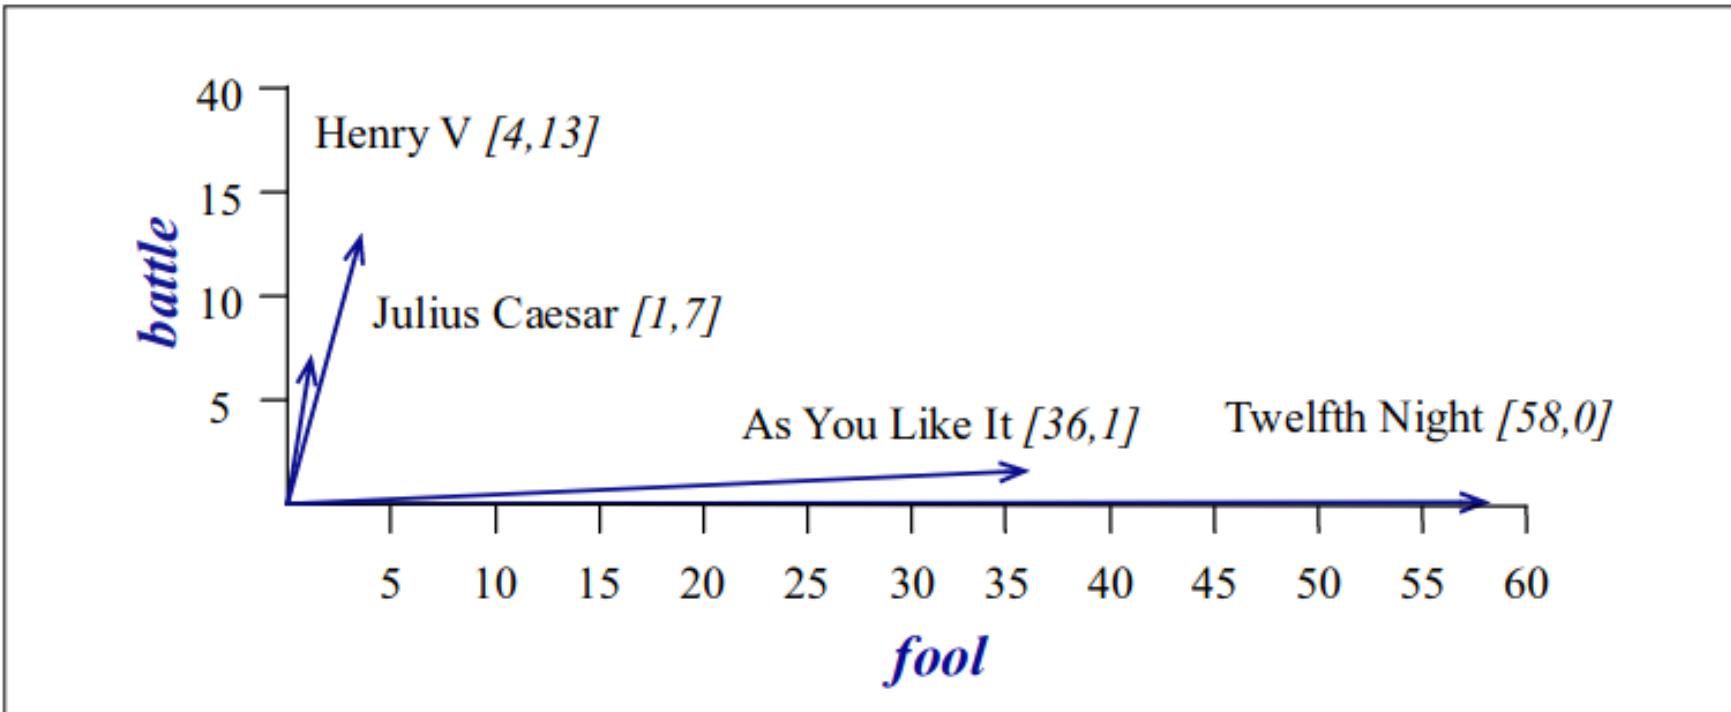


Figure 6.4 A spatial visualization of the document vectors for the four Shakespeare play documents, showing just two of the dimensions, corresponding to the words *battle* and *fool*. The comedies have high values for the *fool* dimension and low values for the *battle* dimension.



Term-Term Matrix

- A term-document matrix represents words as **vectors of their counts across documents**. An alternative is the **term-term matrix** (also called the word-word or term-context matrix), where the columns are labeled by words instead of documents.
 - This matrix is thus of dimensionality $|V| \times |V|$ and each cell records the number of times the row (target) word and the column (context) word co-occur in some context in some training corpus.

| | aardvark | ... | computer | data | result | pie | sugar | ... |
|-------------|----------|-----|----------|------|--------|-----|-------|-----|
| cherry | 0 | ... | 2 | 8 | 9 | 442 | 25 | ... |
| strawberry | 0 | ... | 0 | 0 | 1 | 60 | 19 | ... |
| digital | 0 | ... | 1670 | 1683 | 85 | 5 | 4 | ... |
| information | 0 | ... | 3325 | 3982 | 378 | 5 | 13 | ... |

Figure 6.6 Co-occurrence vectors for four words in the Wikipedia corpus, showing six of the dimensions (hand-picked for pedagogical purposes). The vector for *digital* is outlined in red. Note that a real vector would have vastly more dimensions and thus be much sparser.



Term-Term Matrix

| | aardvark | ... | computer | data | result | pie | sugar | ... |
|-------------|----------|-----|----------|------|--------|-----|-------|-----|
| cherry | 0 | ... | 2 | 8 | 9 | 442 | 25 | ... |
| strawberry | 0 | ... | 0 | 0 | 1 | 60 | 19 | ... |
| digital | 0 | ... | 1670 | 1683 | 85 | 5 | 4 | ... |
| information | 0 | ... | 3325 | 3982 | 378 | 5 | 13 | ... |

Figure 6.6 Co-occurrence vectors for four words in the Wikipedia corpus, showing six of the dimensions (hand-picked for pedagogical purposes). The vector for *digital* is outlined in red. Note that a real vector would have vastly more dimensions and thus be much sparser.

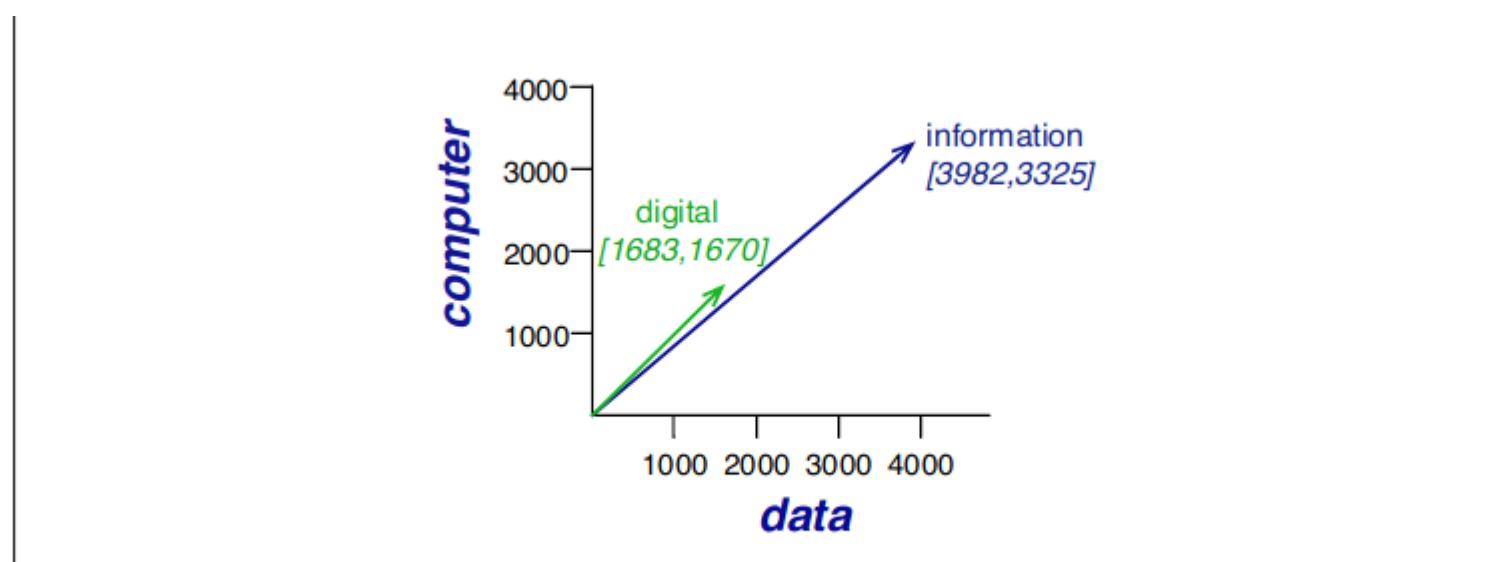
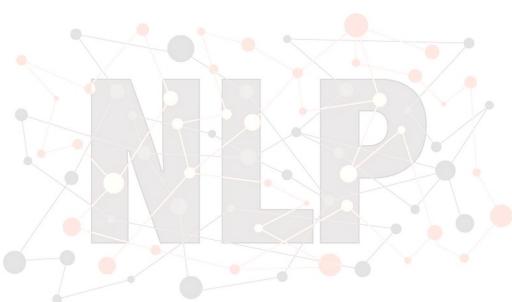


Figure 6.7 A spatial visualization of word vectors for *digital* and *information*, showing just two of the dimensions, corresponding to the words *data* and *computer*.



Cosine for measuring similarity

- To measure similarity between two target words v and w , we need a metric that takes two vectors and gives a measure of their similarity
 - same dimensionality needed
 - either both with words as dimensions, hence of length $|V|$,
 - or both with documents as dimensions, of length $|D|$
- By far the most common similarity metric is the **cosine** of the angle between the vectors

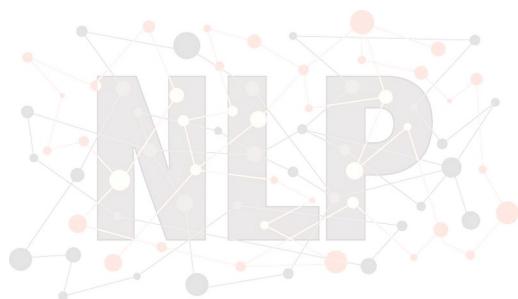


Vector Dot Product

- The dot product acts as a similarity metric because it will tend to be high just when the two vectors have large values in the same dimensions

$$\text{dot product}(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N \quad (6.7)$$

- This raw dot product, however, has a **problem as a similarity metric**: it favors long vectors.
 - higher if a vector is longer
 - higher for frequent words



Normalized dot product

- Same as the cosine of the angle between the two vectors following from the definition of the dot product between two vectors \mathbf{a} and \mathbf{b}

$$\begin{aligned}\mathbf{a} \cdot \mathbf{b} &= |\mathbf{a}| |\mathbf{b}| \cos \theta \\ \frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}| |\mathbf{b}|} &= \cos \theta\end{aligned}$$

- The cosine similarity metric between two vectors \mathbf{v} and \mathbf{w} thus can be computed as:

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$



Cosine Similarity

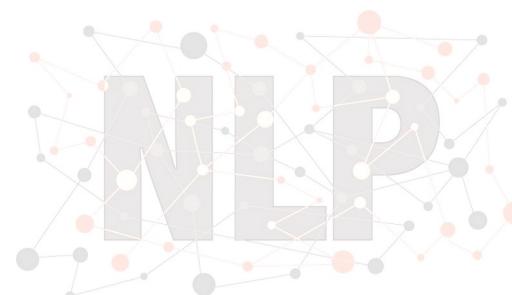
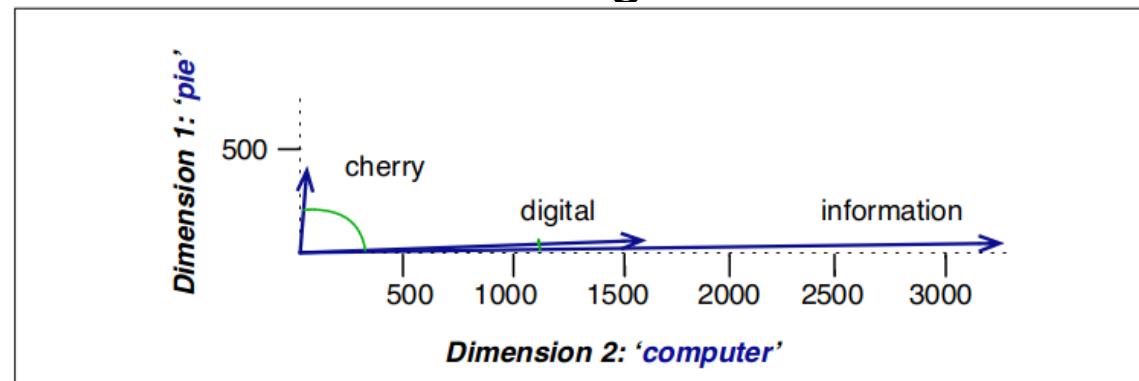
$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

| | pie | data | computer |
|-------------|-----|------|----------|
| cherry | 442 | 8 | 2 |
| digital | 5 | 1683 | 1670 |
| information | 5 | 3982 | 3325 |

$$\cos(\text{cherry}, \text{information}) = \frac{442 * 5 + 8 * 3982 + 2 * 3325}{\sqrt{442^2 + 8^2 + 2^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .018$$

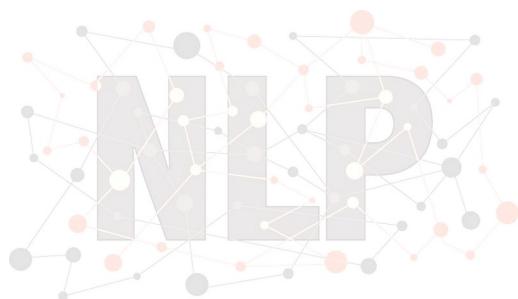
$$\cos(\text{digital}, \text{information}) = \frac{5 * 5 + 1683 * 3982 + 1670 * 3325}{\sqrt{5^2 + 1683^2 + 1670^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .996$$

- When two vectors are more similar, the cosine is larger but the angle is smaller; the cosine has its maximum (1) when the angle between two vectors is smallest (0°); the cosine of all other angles is less than 1



TF-IDF: Weighing terms in the vector

- Words that occur nearby frequently (maybe **pie** nearby **cherry**) are more important than words that only appear once or twice.
- Yet words that are too frequent, like **the** or **good**— are less important
 - How can we balance these two conflicting constraints?



TF-IDF

- Term Frequency — Inverse Document Frequency (TF-IDF)
 - A technique for converting text into **finite length vectors**
 - Gives insights about the less relevant and more relevant words in a document
 - The **importance of a word in the text is of great significance** in information retrieval



Term Frequency

- It is a measure of the frequency of a word (w) in a document (d).
 - TF is defined as the **ratio of a word's occurrence in a document to the total number of words in a document.**
 - The denominator term in the formula is to normalize since all the corpus documents are of different lengths.

$$TF(w, d) = \frac{\text{occurrences of } w \text{ in document } d}{\text{total number of words in document } d}$$



Term Frequency (TF)

- The initial step is to make a vocabulary of unique words and calculate TF for each document.
- TF will be more for words that frequently appear in a document and less for rare words in a document.

| Documents | Text | Total number of words in a document |
|-----------|--|-------------------------------------|
| A | Jupiter is the largest planet | 5 |
| B | Mars is the fourth planet from the sun | 8 |

| Words | TF (for A) | TF (for B) |
|---------|------------|------------|
| Jupiter | 1/5 | 0 |
| Is | 1/5 | 1/8 |
| The | 1/5 | 2/8 |
| largest | 1/5 | 0 |
| Planet | 1/5 | 1/8 |
| Mars | 0 | 1/8 |
| Fourth | 0 | 1/8 |
| From | 0 | 1/8 |
| Sun | 0 | 1/8 |



Inverse Document Frequency (IDF)

- It is the measure of the importance of a word.
 - Term frequency (TF) does not consider the importance of words
 - Some words such as 'of', 'and', etc. can be most frequently present but are of little significance
 - IDF provides weightage to each word based on its frequency in the corpus **D**



Inverse Document Frequency (IDF)

- IDF of a word (w) is defined as
 - $\ln = \log_e$

$$IDF(w, D) = \ln\left(\frac{\text{Total number of documents } (N) \text{ in corpus } D}{\text{number of documents containing } w}\right)$$



Inverse Document Frequency (IDF)

- In our example, since we have two documents in the corpus, N=2.

| Words | TF (for A) | TF (for B) | IDF |
|---------|------------|------------|-------------------|
| Jupiter | 1/5 | 0 | $\ln(2/1) = 0.69$ |
| Is | 1/5 | 1/8 | $\ln(2/2) = 0$ |
| The | 1/5 | 2/8 | $\ln(2/2) = 0$ |
| largest | 1/5 | 0 | $\ln(2/1) = 0.69$ |
| Planet | 1/5 | 1/8 | $\ln(2/2) = 0$ |
| Mars | 0 | 1/8 | $\ln(2/1) = 0.69$ |
| Fourth | 0 | 1/8 | $\ln(2/1) = 0.69$ |
| From | 0 | 1/8 | $\ln(2/1) = 0.69$ |
| Sun | 0 | 1/8 | $\ln(2/1) = 0.69$ |



TF-IDF

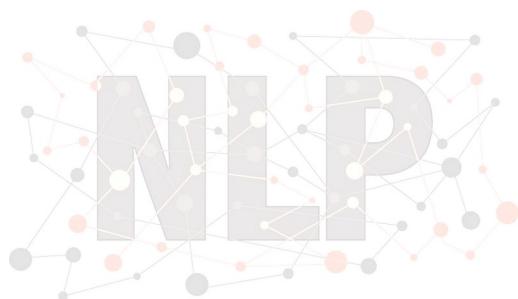
- It is the product of TF and IDF.
 - TFIDF gives more weightage to the word that is **rare in the corpus** (all the documents).
 - TFIDF provides more importance to the word that is more frequent in the document.

$$TFIDF(w, d, D) = TF(w, d) * IDF(w, D)$$



Why Ln in the IDF?

- TFIDF is the product of TF with IDF.
 - Since TF values lie between 0 and 1,
 - Not using **ln** can result in high IDF for some words, thereby dominating the TFIDF. We don't want that, and therefore
- We use **ln** so that IDF should not completely dominate the TFIDF.



TF-IDF

| Words | TF (for A) | TF (for B) | IDF | TFIDF (A) | TFIDF (B) |
|---------|------------|------------|-------------------|-----------|-----------|
| Jupiter | 1/5 | 0 | $\ln(2/1) = 0.69$ | 0.138 | 0 |
| Is | 1/5 | 1/8 | $\ln(2/2) = 0$ | 0 | 0 |
| The | 1/5 | 2/8 | $\ln(2/2) = 0$ | 0 | 0 |
| largest | 1/5 | 0 | $\ln(2/1) = 0.69$ | 0.138 | 0 |
| Planet | 1/5 | 1/8 | $\ln(2/2) = 0$ | 0.138 | 0 |
| Mars | 0 | 1/8 | $\ln(2/1) = 0.69$ | 0 | 0.086 |
| Fourth | 0 | 1/8 | $\ln(2/1) = 0.69$ | 0 | 0.086 |
| From | 0 | 1/8 | $\ln(2/1) = 0.69$ | 0 | 0.086 |
| Sun | 0 | 1/8 | $\ln(2/1) = 0.69$ | 0 | 0.086 |



Disadvantage of TFIDF

- It is **unable to capture the semantics**
 - For example, funny and humorous are synonyms, but TFIDF does not capture that.
- Moreover, TFIDF can be **computationally expensive** if the vocabulary is vast.



Pointwise mutual information

- It is a measure of **how often two events x and y occur**, compared with what we would expect if they were independent

$$I(x, y) = \log_2 \frac{P(x, y)}{P(x)P(y)} \quad (6.16)$$

The pointwise mutual information between a target word w and a context word c (Church and Hanks 1989, Church and Hanks 1990) is then defined as:

$$\text{PMI}(w, c) = \log_2 \frac{P(w, c)}{P(w)P(c)} \quad (6.17)$$

Read more – yourself. Chapter 6 – PMI, PPMI



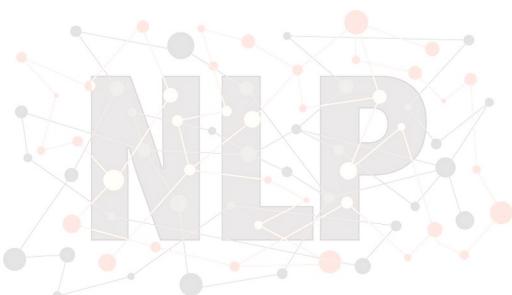
Text Embedding

Word2vec



Word2vec

- Unlike the vectors we've seen so far, **embeddings are short**, with number of dimensions **d** ranging from **50-1000**, rather than the much **larger vocabulary size |V|** or **number of documents D** we've seen
 - These **d** dimensions don't have a clear interpretation
 - The vectors are dense: instead of vector entries being sparse, mostly-zero counts or functions of counts, the values will be real-valued numbers that can be negative



Word2vec: Skip-grams

- Instead of entire documents, Word2vec uses words a few positions away from each center word. The pairs of center word/context word are called “skip-grams”

“It was **a bright cold day in April, and** the clocks were striking”

Center word: red

Context words: blue

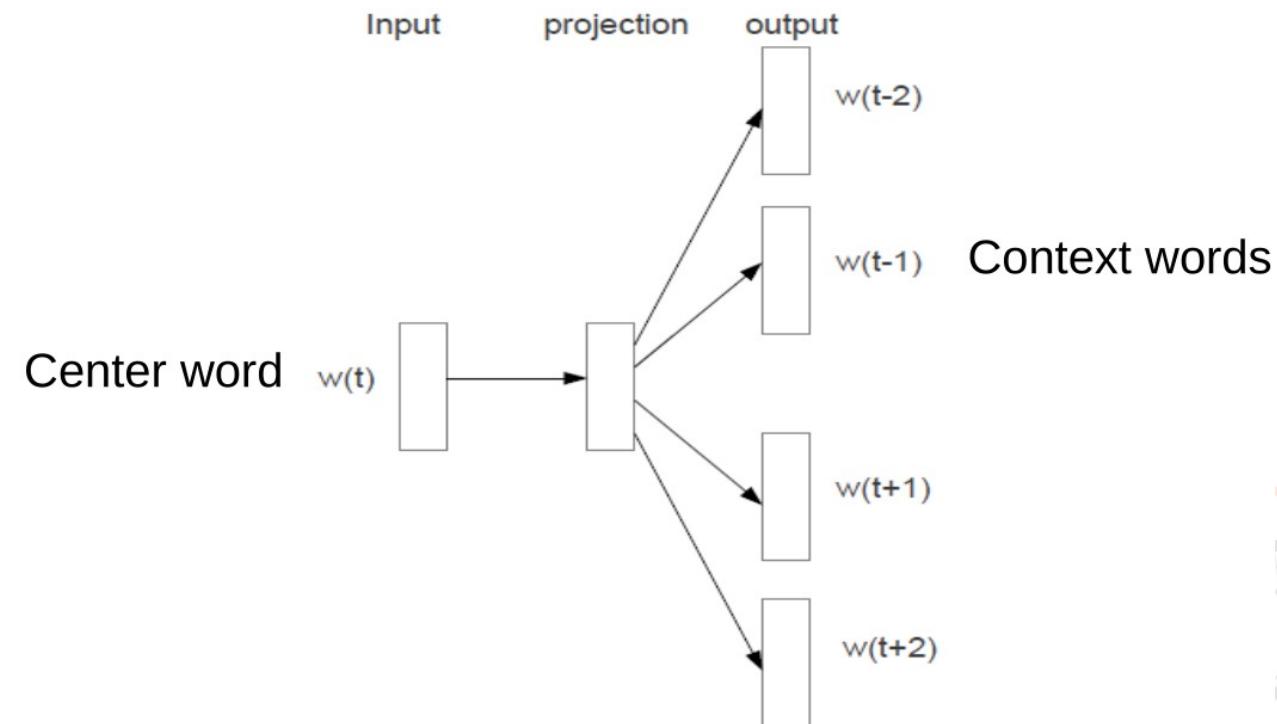
- Word2vec considers all words as center words, and all their context words



Word2vec: Skip-grams

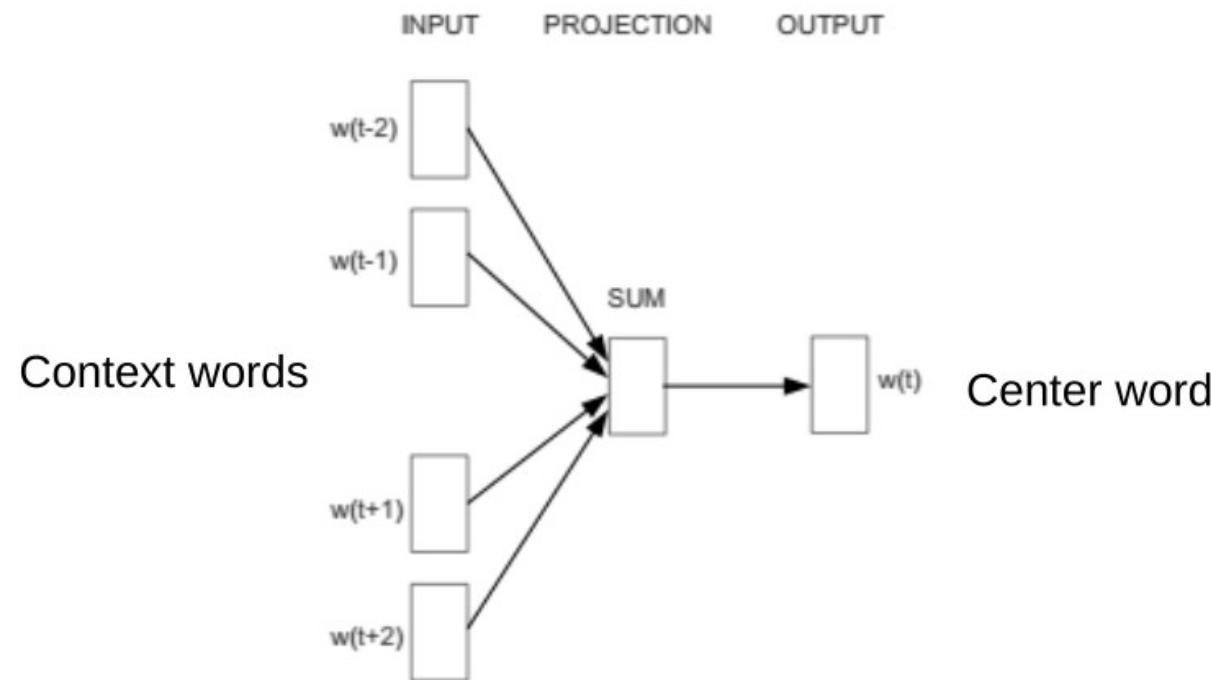
- The pairs of center word/context word are called “**skip-grams**”
- Typical distances are 3-5 word positions.

Skip-gram model



Word2vec: CBOW

- Models can also predict center word from context, CBOW model. Generally, skip-gram performs better



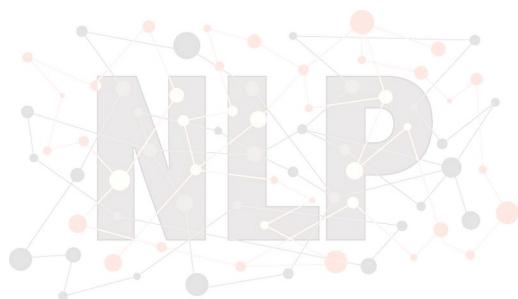
Skip-gram model

- The goal of the skip-gram model is to **predict the context words given a target word**.
- Specifically, it maximizes the probability of the surrounding words within a certain window size based on the current word.
- Skip-gram tends to work better when data is sparse and can learn high-quality embeddings even from large corpora



Skip-gram – Objective Function

- Given a corpus of words $\{w_1, w_2, \dots, w_T\}$
 - the skip-gram model aims to maximize the average log probability of the context words
 w_{t-j}, \dots, w_{t+j}
 - given the target word
 w_t (excluding w_t itself)
 - over all positions t in the corpus.



Skip-gram – Objective Function

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j} | w_t)$$

where

- w_t is the target word at position t,
- w_{t+j} are the context words surrounding w_t ,
- c is the window size, determining how many words before and after the target word are considered as context,
- T is the total number of words in the corpus.

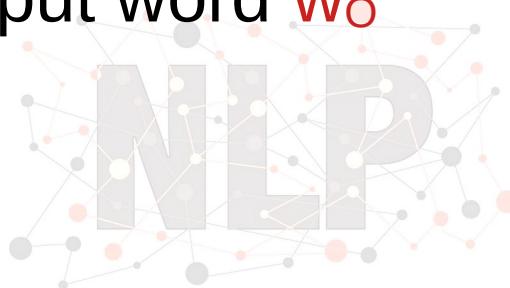


Skip-gram – Objective Function

- The conditional probability $P(w_{t+j} | w_t)$ is typically modeled using a softmax function over the entire vocabulary V , parameterized by the word embeddings:

$$P(w_O | w_I) = \frac{\exp(v_{w_O}^\top v_{w_I})}{\sum_{w \in V} \exp(v_w^\top v_{w_I})}$$

- Where:
 - v_{w_I} is the vector representation (embedding) of the input word w_I ,
 - v_{w_O} is the vector representation (embedding) of the output word w_O
 - V is the vocabulary



Word2Vec

Word2vec - DEMO



Thank you

